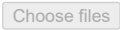


## Upload the Dataset

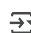
```
from google.colab import files
uploaded = files.upload()
```

 No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable

## Load the Dataset

```
import pandas as pd

# Load dataset
df = pd.read_csv('/content/raw_sales (1).csv')
df.head()
```




|   | datesold            | postcode | price  | propertyType | bedrooms |
|---|---------------------|----------|--------|--------------|----------|
| 0 | 2007-02-07 00:00:00 | 2607     | 525000 | house        | 4        |
| 1 | 2007-02-27 00:00:00 | 2906     | 290000 | house        | 3        |
| 2 | 2007-03-07 00:00:00 | 2905     | 328000 | house        | 3        |
| 3 | 2007-03-09 00:00:00 | 2905     | 380000 | house        | 4        |
| 4 | 2007-03-21 00:00:00 | 2906     | 310000 | house        | 3        |

## Data Exploration

```
# Dataset Info
df.info()

# Summary Statistics
df.describe()

# First few records
df.head()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29580 entries, 0 to 29579
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   datesold        29580 non-null  object
1   postcode        29580 non-null  int64
2   price           29580 non-null  int64
3   propertyType    29580 non-null  object
4   bedrooms        29580 non-null  int64
dtypes: int64(3), object(2)
memory usage: 1.1+ MB
```

|   | datesold            | postcode | price  | propertyType | bedrooms |
|---|---------------------|----------|--------|--------------|----------|
| 0 | 2007-02-07 00:00:00 | 2607     | 525000 | house        | 4        |
| 1 | 2007-02-27 00:00:00 | 2906     | 290000 | house        | 3        |
| 2 | 2007-03-07 00:00:00 | 2905     | 328000 | house        | 3        |
| 3 | 2007-03-09 00:00:00 | 2905     | 380000 | house        | 4        |
| 4 | 2007-03-21 00:00:00 | 2906     | 310000 | house        | 3        |

## Check for Missing Values and Duplicates

```
import pandas as pd

# Example: Create a sample DataFrame or load one
# Option 1: Create manually
data = {'A': [1, 2, 2, None], 'B': [4, None, 4, 4]}
df = pd.DataFrame(data)





# Option 2: Load from a file (e.g., CSV)
# df = pd.read_csv('your_file.csv')
```

## Visualize a Few Features

```
import matplotlib.pyplot as plt
import seaborn as sns

# Histogram of numerical columns
df.hist(bins=30, figsize=(12, 10))
plt.tight_layout()
plt.show()

# Correlation heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```



## Identify Target and Features

```
import os
print("Current working directory:", os.getcwd())
print("Files in current directory:", os.listdir())
```

```
↗ Current working directory: /content
Files in current directory: ['.config', 'raw_sales (1).csv', 'sample_data']
```

## Convert Categorical Columns to Numerical

```
import os

print("Current working directory:", os.getcwd())
print("Files in this directory:", os.listdir())
```

```
↗ Current working directory: /content
Files in this directory: ['.config', 'raw_sales (1).csv', 'sample_data']
```

## One-Hot Encoding

```
import os

# Check where Python is looking
print("Current working directory:", os.getcwd())

# List all files in that directory
print("Files in this directory:", os.listdir())
```

```
↗ Current working directory: /content
Files in this directory: ['.config', 'raw_sales (1).csv', 'sample_data']
```

## Feature Scaling

```
import pandas as pd
from sklearn.preprocessing import StandardScaler

# Load the dataset
df = pd.read_csv('/content/raw_sales (1).csv') # Make sure this file exists

# Define X
X = df.copy() # Or df.drop('target_column', axis=1)

# Identify categorical columns
categorical_cols = X.select_dtypes(include=['object']).columns.tolist()

# One-hot encode
X_encoded = pd.get_dummies(X, columns=categorical_cols, drop_first=True)

# Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_encoded)

# (Optional) Convert scaled array back to DataFrame
X_scaled_df = pd.DataFrame(X_scaled, columns=X_encoded.columns)
print(X_scaled_df.head())
```

```
↗
```

|   | postcode  | price     | bedrooms  | datesold_2007-02-27 00:00:00 \ |
|---|-----------|-----------|-----------|--------------------------------|
| 0 | -0.840063 | -0.300800 | 0.788251  | -0.005814                      |
| 1 | 1.197904  | -1.135011 | -0.262987 | 171.985465                     |
| 2 | 1.191088  | -1.000118 | -0.262987 | -0.005814                      |
| 3 | 1.191088  | -0.815526 | 0.788251  | -0.005814                      |
| 4 | 1.197904  | -1.064015 | -0.262987 | -0.005814                      |

|   | datesold_2007-03-07 00:00:00 | datesold_2007-03-09 00:00:00 \ |
|---|------------------------------|--------------------------------|
| 0 | -0.005814                    | -0.005814                      |
| 1 | -0.005814                    | -0.005814                      |

```

2          171.985465          -0.005814
3          -0.005814          171.985465
4          -0.005814          -0.005814

    datesold_2007-03-21 00:00:00 datesold_2007-04-04 00:00:00 \
0          -0.005814          -0.005814
1          -0.005814          -0.005814
2          -0.005814          -0.005814
3          -0.005814          -0.005814
4          171.985465          -0.005814

    datesold_2007-04-24 00:00:00 datesold_2007-04-30 00:00:00 ... \
0          -0.005814          -0.005814 ...
1          -0.005814          -0.005814 ...
2          -0.005814          -0.005814 ...
3          -0.005814          -0.005814 ...
4          -0.005814          -0.005814 ...

    datesold_2019-07-18 00:00:00 datesold_2019-07-19 00:00:00 \
0          -0.017446          -0.01839
1          -0.017446          -0.01839
2          -0.017446          -0.01839
3          -0.017446          -0.01839
4          -0.017446          -0.01839

    datesold_2019-07-20 00:00:00 datesold_2019-07-22 00:00:00 \
0          -0.015385          -0.015385
1          -0.015385          -0.015385
2          -0.015385          -0.015385
3          -0.015385          -0.015385
4          -0.015385          -0.015385

    datesold_2019-07-23 00:00:00 datesold_2019-07-24 00:00:00 \
0          -0.022525          -0.011629
1          -0.022525          -0.011629
2          -0.022525          -0.011629
3          -0.022525          -0.011629
4          -0.022525          -0.011629

    datesold_2019-07-25 00:00:00 datesold_2019-07-26 00:00:00 \
0          -0.022525          -0.014244
1          -0.022525          -0.014244
2          -0.022525          -0.014244
3          -0.022525          -0.014244
4          -0.022525          -0.014244

    datesold_2019-07-27 00:00:00 propertyType_unit
0          -0.010071          -0.452537

```

## Train-Test Split

```

# Target and features
target = 'price'
features = ['datesold', 'postcode', 'propertyType', 'bedrooms']

X = df[features]
y = df[target]

```

## Model Building

```

import os
print(os.getcwd())

```

 /content

## Evaluation

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

# Generate mock data
np.random.seed(42)
n_samples = 500
df = pd.DataFrame({
    'datesold': pd.date_range(start='2018-01-01', periods=n_samples, freq='D'),
    'postcode': np.random.randint(1000, 9999, size=n_samples),
    'propertyType': np.random.choice(['House', 'Unit', 'Townhouse'], size=n_samples),
    'bedrooms': np.random.randint(1, 5, size=n_samples),

```

```

    'price': np.random.randint(100000, 1000000, size=n_samples)
})

# Target and features
target = 'price'
features = ['datesold', 'postcode', 'propertyType', 'bedrooms']

X = df[features]
y = df[target]

# Convert 'datesold' to datetime and extract useful features
X['datesold'] = pd.to_datetime(X['datesold'])
X['year'] = X['datesold'].dt.year
X['month'] = X['datesold'].dt.month
X = X.drop('datesold', axis=1)

# One-hot encode categorical column
X = pd.get_dummies(X, columns=['propertyType'], drop_first=True)

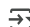
# Scale features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42
)

# Train model
model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train, y_train)

# Evaluate
y_pred = model.predict(X_test)
print("R² Score:", r2_score(y_test, y_pred))
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))

```

 <ipython-input-14-2d032475104e>:27: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus)  
 X['datesold'] = pd.to\_datetime(X['datesold'])  
 R² Score: -0.22806619582534382  
 Mean Squared Error: 75459008443.73859

## Make Predictions from New Input

```

# Example dictionary (adjust keys to match your actual features)
new_data = {
    'Bedrooms': 3,
    'Bathrooms': 2,
    'SqFt': 1500,
    'Location': 'Downtown', # Example categorical
    # Add other features as needed...
}
new_df = pd.DataFrame([new_data])

```

## Convert to DataFrame and Encode

```

import pandas as pd
import numpy as np
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score

# Create mock data
np.random.seed(42)
n_samples = 300
df = pd.DataFrame({
    'datesold': pd.date_range(start='2021-01-01', periods=n_samples, freq='D'),
    'postcode': np.random.randint(1000, 9999, n_samples),
    'propertyType': np.random.choice(['House', 'Unit', 'Townhouse'], size=n_samples),
    'bedrooms': np.random.randint(1, 5, n_samples),
    'price': np.random.randint(150000, 1000000, n_samples)
})

```

```

# Preprocessing date
df['datesold'] = pd.to_datetime(df['datesold'])
df['year'] = df['datesold'].dt.year
df['month'] = df['datesold'].dt.month
df = df.drop(columns=['datesold'])

# Split features and target
X = df.drop(columns=['price'])
y = df['price']

# Define column types
categorical_cols = ['propertyType']
numerical_cols = ['postcode', 'bedrooms', 'year', 'month']

# Preprocessing pipeline
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numerical_cols),
        ('cat', OneHotEncoder(drop='first'), categorical_cols)
    ]
)

# Full pipeline with model
model_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('regressor', RandomForestRegressor(n_estimators=100, random_state=42))
])

# Split and train
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model_pipeline.fit(X_train, y_train)

# Predict and evaluate
y_pred = model_pipeline.predict(X_test)
print("R² Score:", r2_score(y_test, y_pred))
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))

↗ R² Score: -0.18492132175595444
Mean Squared Error: 66698791659.67322

```

### Predict the Final Price

```

import glob
import os

search_path = os.path.expanduser("")
for file in glob.glob(search_path, recursive=True):
    print(file)

```

### Deployment - Building an Interactive App (Gradio)

```

!pip install gradio

import gradio as gr

def predict_price(**kwargs):
    input_df = pd.DataFrame([kwargs])
    input_encoded = pd.get_dummies(input_df)
    input_encoded = input_encoded.reindex(columns=X_encoded.columns, fill_value=0)
    input_scaled = scaler.transform(input_encoded)
    prediction = model.predict(input_scaled)
    return "${:,.2f}".format(prediction[0])

↗ Collecting gradio
  Downloading gradio-5.29.0-py3-none-any.whl.metadata (16 kB)
Collecting aiofiles<25.0,>=22.0 (from gradio)
  Downloading aiofiles-24.1.0-py3-none-any.whl.metadata (10 kB)
Requirement already satisfied: anyio<5.0,>=3.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (4.9.0)
Collecting fastapi<1.0,>=0.115.2 (from gradio)
  Downloading fastapi-0.115.12-py3-none-any.whl.metadata (27 kB)
Collecting ffmpy (from gradio)
  Downloading ffmpy-0.5.0-py3-none-any.whl.metadata (3.0 kB)
Collecting gradio-client==1.10.0 (from gradio)
  Downloading gradio_client-1.10.0-py3-none-any.whl.metadata (7.1 kB)
Collecting groovy~=0.1 (from gradio)
  Downloading groovy-0.1.2-py3-none-any.whl.metadata (6.1 kB)
Requirement already satisfied: httpx>=0.24.1 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.28.1)
Requirement already satisfied: huggingface-hub>=0.28.1 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.30.2)
Requirement already satisfied: jinja2<4.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.1.6)
Requirement already satisfied: markupsafe<4.0,>=2.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.0.2)

```

Requirement already satisfied: numpy<3.0,>=1.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.0.2)  
Requirement already satisfied: orjson~=3.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.10.18)  
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from gradio) (24.2)  
Requirement already satisfied: pandas<3.0,>=1.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.2.2)  
Requirement already satisfied: pillow<12.0,>=8.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (11.2.1)  
Requirement already satisfied: pydantic<2.12,>=2.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.11.4)  
Collecting pydub (from gradio)  
 Downloading pydub-0.25.1-py2.py3-none-any.whl.metadata (1.4 kB)  
Collecting python-multipart>=0.0.18 (from gradio)  
 Downloading python\_multipart-0.0.20-py3-none-any.whl.metadata (1.8 kB)  
Requirement already satisfied: pyyaml<7.0,>=5.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (6.0.2)